Triage Report — Synapse Planners

1. Submission Summary

Synapse Planners propose an AI-driven autonomous inventory management solution for retailers using Microsoft Copilot Studio, Dynamics 365 integration, Power Automate, Azure AI Foundry, Teams, and Purview. The concept focuses on monitoring stock, forecasting demand, detecting anomalies, and triggering replenishment or escalation workflows with human oversight and policy guardrails.

- **Industry:** Retail
- **Use Case:** Retail — **Purchase Planner Approach / Replenishment Copilot**, with some elements of the **Planner exception — Demand Planner Assistant** and partial ambition toward an autonomous control tower
- **Track: C** (Dynamics 365 + Copilot Studio + Power Platform + optional GitHub + Azure) inferred from the stated stack: Dynamics 365, Copilot Studio, Power Automate, Azure AI Foundry
- **Stage:** Submission Review

Cross-cutting checks

Check	Status	Notes
Environment Track declared	Warn	The submission does not explicitly state Track C, but the architecture clearly aligns best with **Track C** through Dynamics 365 + Copilot Studio + Power Platform, with optional Azure AI Foundry.
Industry Use Case aligned	Pass	Strong fit to **Retail — Autonomous Inventory Management**. Most closely matches the defined **Purchase Planner Approach — Replenishment Copilot** because it watches exceptions and proposes/executes replenishment via connectors with approvals.
Azure WAF aligned	Warn	Azure AI Foundry is included, but Azure WAF coverage is only partial. Performance intent is described, but reliability, cost optimisation, and operational excellence controls are not concretely defined.
Azure WAF AI Pillar aligned	Warn	Agent design and basic human oversight are described well. Responsible AI and guardrails are partially covered, but model lifecycle, grounding strategy, AI observability, and evaluation processes are only lightly addressed.
Power Platform WAF aligned	Warn	Copilot Studio and Power Automate are appropriate and well aligned functionally, but the submission lacks details on environment strategy, ALM, solution packaging, resiliency, and experience optimisation beyond Teams/adaptive cards.
Microsoft Security & Monitoring	Warn	Purview and auditability are mentioned, but identity, RBAC, monitoring, alerting, secrets handling, DLP, and end-to-end operational monitoring are not sufficiently specified.

2. Criteria Scoring

Architecture

Dimension	Score	Justification

Clarity & Flow	**8/10**	The end-to-end design is clearly articulated through architecture layers and a 6-step workflow from monitor → forecast → detect → policy check → execute → learn. This scores well on **Clear Architecture Flow**, though some flows remain conceptual rather than system-sequence-level detailed.
Platform Fit	**8/10**	The chosen stack is a strong match for the problem and closely aligns with **Track C**: Copilot Studio for agent interaction, Dynamics 365 for ERP context, Power Automate for orchestration, Azure AI Foundry for forecasting, and Teams for approvals. Score is reduced because the rationale for when to use Copilot Studio vs. Azure AI Foundry vs. MCP is not fully differentiated.
Data / Governance / Security	**6/10**	The submission identifies relevant data sources and mentions Purview auditability, policy guardrails, confidence thresholds, and approval rules. However, against **Data & Governance** and **Security**, it lacks specifics on data ownership, data storage patterns, identity model, access controls, DLP, retention, sensitive supplier/commercial data handling, and content safety controls.
Reliability / Performance / Cost	**5/10**	The design implies autonomous and escalated paths and includes some fallback through human approval, which helps **Reliability**. But there is little on retries, failure recovery, degraded modes, API outage handling, latency expectations, throughput, or **Cost Awareness** for AI inference, forecasting cadence, and integration calls.
Scalability / Integration / Provisioning	**6/10**	Integrations are identified clearly across ERP, POS, supplier APIs, e-commerce, Teams, and Power Automate, which supports the **Integration** criterion. Scalability is plausible, but provisioning readiness is weak: no concrete deployment topology, environments, connectors list, IaC, or implementation boundaries are defined.

Architecture Subtotal: 33/50

Use Case

Dimension	Score	Justification
Value & KPI Impact	**8/10**	This addresses a significant retail problem with broad user and business value: stockouts, overstock, markdowns, inventory turns, labor savings, and supplier coordination. It performs strongly on **Problem Significance** and **User Impact**. Score is slightly reduced because KPIs are described qualitatively, with few target metrics or baselines.
Innovation	**7/10**	The combination of autonomous replenishment, supplier scoring, policy-based approvals, external demand signals, and Teams-based human-in-the-loop is thoughtful and relevant. However, much of it reflects established inventory AI patterns rather than a distinctly novel mechanism beyond current market direction.
AI Fit	**8/10**	AI is well suited to forecasting, anomaly detection, and adaptive replenishment where static rules are insufficient. The submission also includes some **Responsible Design** through policy guardrails, confidence thresholds, approvals, and auditability. The score is reduced because fairness, explainability, hallucination handling, and explicit responsible AI controls are not deeply covered.
UX Simplicity	**7/10**	The proposed UX is intuitive: Teams/Copilot chat, adaptive cards, natural language queries, and clear approve/override loops. This meets the **User Experience** criterion reasonably well. More concrete user journeys for planners, store managers, and procurement users would strengthen confidence.

Build–Scale–Reuse	**6/10**	The scope is credible at prototype level, especially with Copilot Studio and Power Automate accelerating delivery, and there is reuse potential through MCP/connectors and reusable approval/agent patterns. However, the end-to-end ambition across forecasting, supplier APIs, auto-POs, stock rebalancing, markdown triggers, and learning loops may be too broad for hackathon execution without narrowing.

Use Case Subtotal: 36/50

Total

Total Score: 69/100

---

3. Strengths

- Strong alignment to the **Retail — Autonomous Inventory Management** scenario, especially the **Purchase Planner Approach — Replenishment Copilot** criterion, with a clear business problem and credible agent role.
- Good **Clarity & Flow**: the 6-step decision loop is easy to follow and makes the autonomous vs. human-escalated decision path understandable.
- Solid **Platform Fit** for Track C: Copilot Studio, Dynamics 365, Power Automate, Teams, and Azure AI Foundry are all relevant choices for this workflow.
- Good **AI Fit**: forecasting, anomaly detection, and supplier decision support are areas where AI adds meaningful value beyond simple threshold rules.
- Thoughtful governance foundations through policy guardrails, confidence thresholds, high-value order approvals, and Purview-based audit trail.
- Clear articulation of business value areas including stockout reduction, fewer markdowns, higher inventory turns, and labor savings, supporting the **Value & KPI Impact** criterion.

---

4. Gaps & Risks

- **Environment Track** is not explicitly declared, which weakens submission completeness and creates ambiguity against the track-alignment check.
- Under **Data / Governance / Security**, the architecture does not define where forecast data, event data, supplier data, and action logs are stored, who owns them, or how access is controlled.
- Under **Microsoft Security & Monitoring**, there is no concrete mention of Entra ID/RBAC, secrets management, API authentication, DLP policies, logging, alerting, or operational dashboards.
- Under **Reliability / Performance / Cost**, failure handling is underdeveloped: no retry strategy, no supplier API downtime path, no manual fallback detail beyond approvals, and no AI/token/call cost controls.
- Under **Scalability / Integration / Provisioning**, the submission is not yet provisioning-ready: no environment design, no deployment approach, no connector inventory, and no description of how MCP/connectors are packaged or governed.
- Under **AI Fit / Responsible Design**, explainability for forecast-driven reorder decisions is not explicit; this could reduce trust for planners approving or auditing decisions.
- Under **Build–Scale–Reuse**, the scope may be too large for a hackathon if all features are attempted at once; the current design spans forecasting, automation, supplier orchestration, transfers, markdowns, and feedback learning.
- The use case mentions external signals like weather and social trends, but does not explain data quality thresholds, grounding, or how noisy signals are prevented from causing bad replenishment actions.

---

5. Triage Verdict

Good — 69/100

---

6. Top Actionable Improvements

- **Explicitly declare Track C and tighten component-to-track mapping** so the submission clearly shows how Dynamics 365, Copilot Studio, Power Automate, Teams, and optional Azure AI Foundry each map to the chosen hackathon environment and use case criteria.
- **Strengthen Data / Governance / Security** by adding a concrete data architecture: system of record, where forecasts and action logs live, identity model, RBAC, DLP, retention, supplier data confidentiality controls, and approval/audit boundaries.
- **Add Reliability / Performance / Cost controls** including retry/fallback logic for ERP and supplier API failures, escalation when models are low confidence, expected response times for planner workflows, and cost controls for forecast frequency and AI inference usage.
- **Make the solution provisioning-ready** by listing required services, connectors, environments, deployment approach, and whether any reusable MCP/connector assets will be produced, directly improving the **Scalability / Integration / Provisioning** criterion.
- **Quantify KPI impact and narrow MVP scope** by defining 3–5 measurable target outcomes such as stockout reduction %, forecast accuracy uplift, planner hours saved, markdown reduction %, and then prioritising a hackathon MVP around one core closed loop: demand exception → approval → PO creation.

Important Notice — AI-Assisted Evaluation
This report has been produced using an AI-powered triage system developed by Capgemini as part of the Microsoft Agentic Industry Hackathon 2026 evaluation process. While every effort has been made to ensure consistency and quality, the scores and assessments contained herein are indicative only and should not be treated as definitive or final judgements. AI-generated evaluations may contain inaccuracies, omissions, or contextual misinterpretations.

This report is intended as a supporting tool to assist Capgemini and Microsoft reviewers and also participating teams, judges, and Partner Solution Architects (PSAs) in identifying teams that are currently active and assessing relative submission strength, with the aim of informing — not replacing — human decision-making on stage progression. All final advancement decisions will be made by qualified human reviewers from Capgemini and Microsoft, who will use this report as one input alongside their own professional assessment.

Support & Enquiries: Clemens Reijnen <clemens.reijnen@sogeti.com> | Logan Talbot <logan.talbot@capgemini.com> | v0.0.2

Capgemini